




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Parallel I/O on Piz Daint

User Lab Day 2018

Samuel Omlin, CSCS

September 11<sup>th</sup> 2018

# Outline

- Introduction to I/O on Piz Daint
  - Scratch – a Lustre file system
- Parallel I/O? – Common approaches
  - File-per-process
  - Shared file
  - Recommendations for Piz Daint
- General recommendations for any I/O approach on Piz Daint
- Conclusions



*CSCS office building in Lugano*



**CSCS**

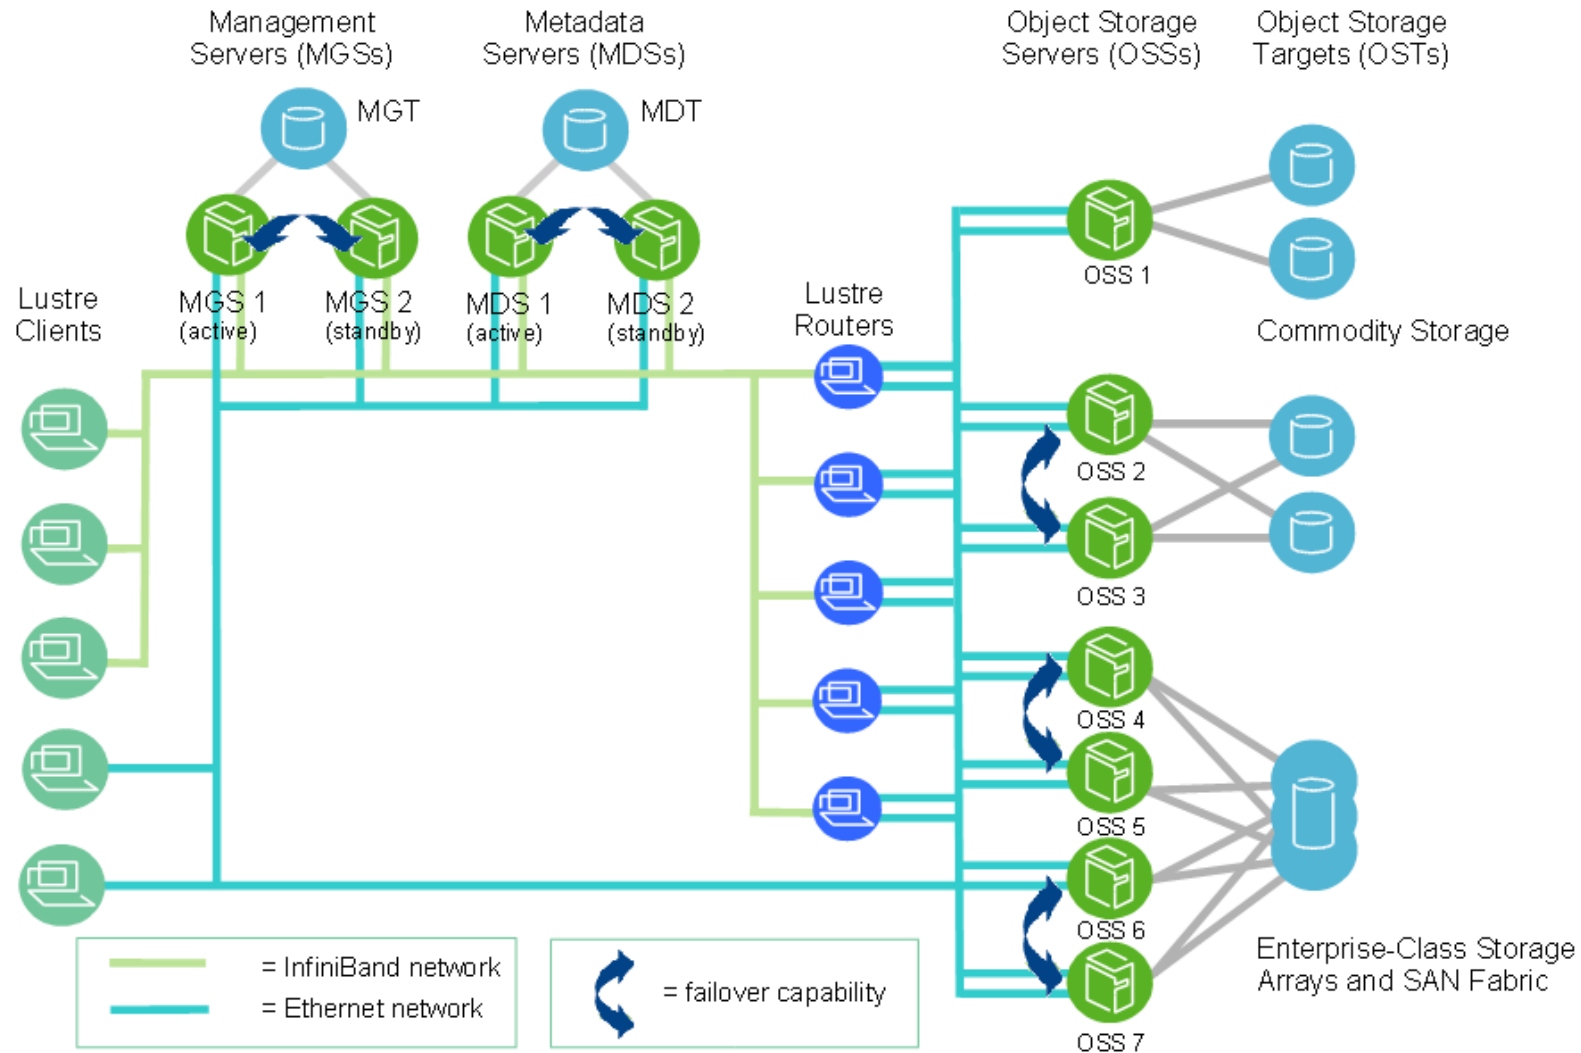
Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

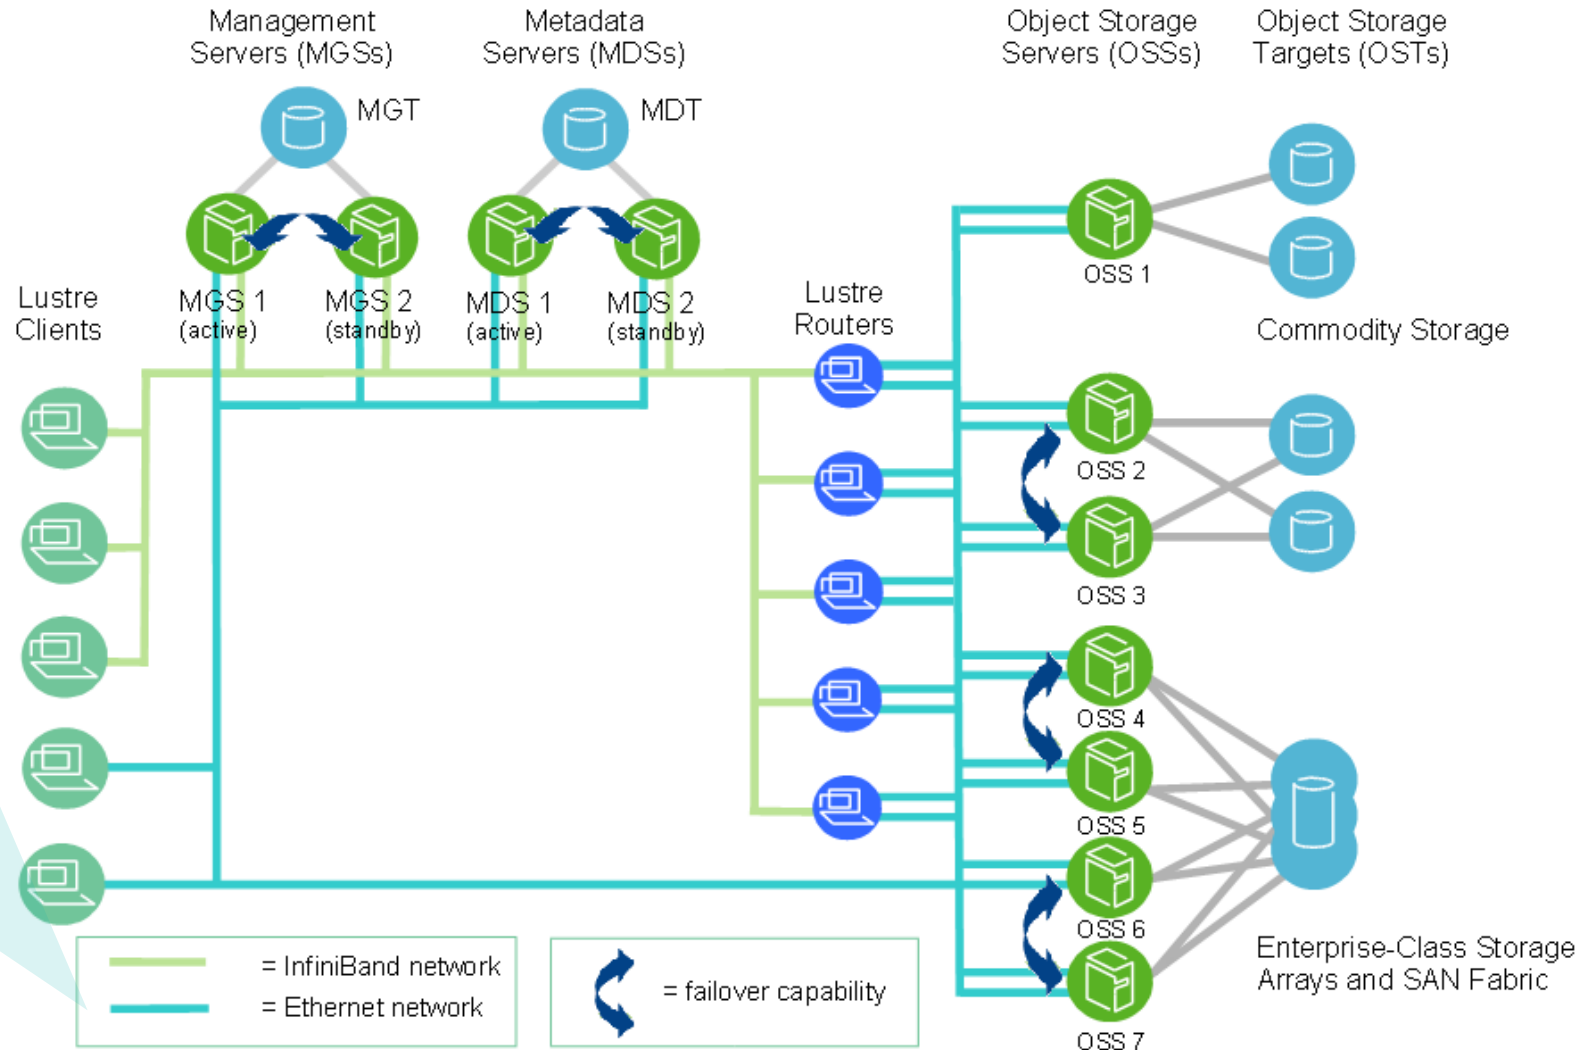
# Introduction to I/O on Piz Daint

---

# Scratch – a Lustre file system



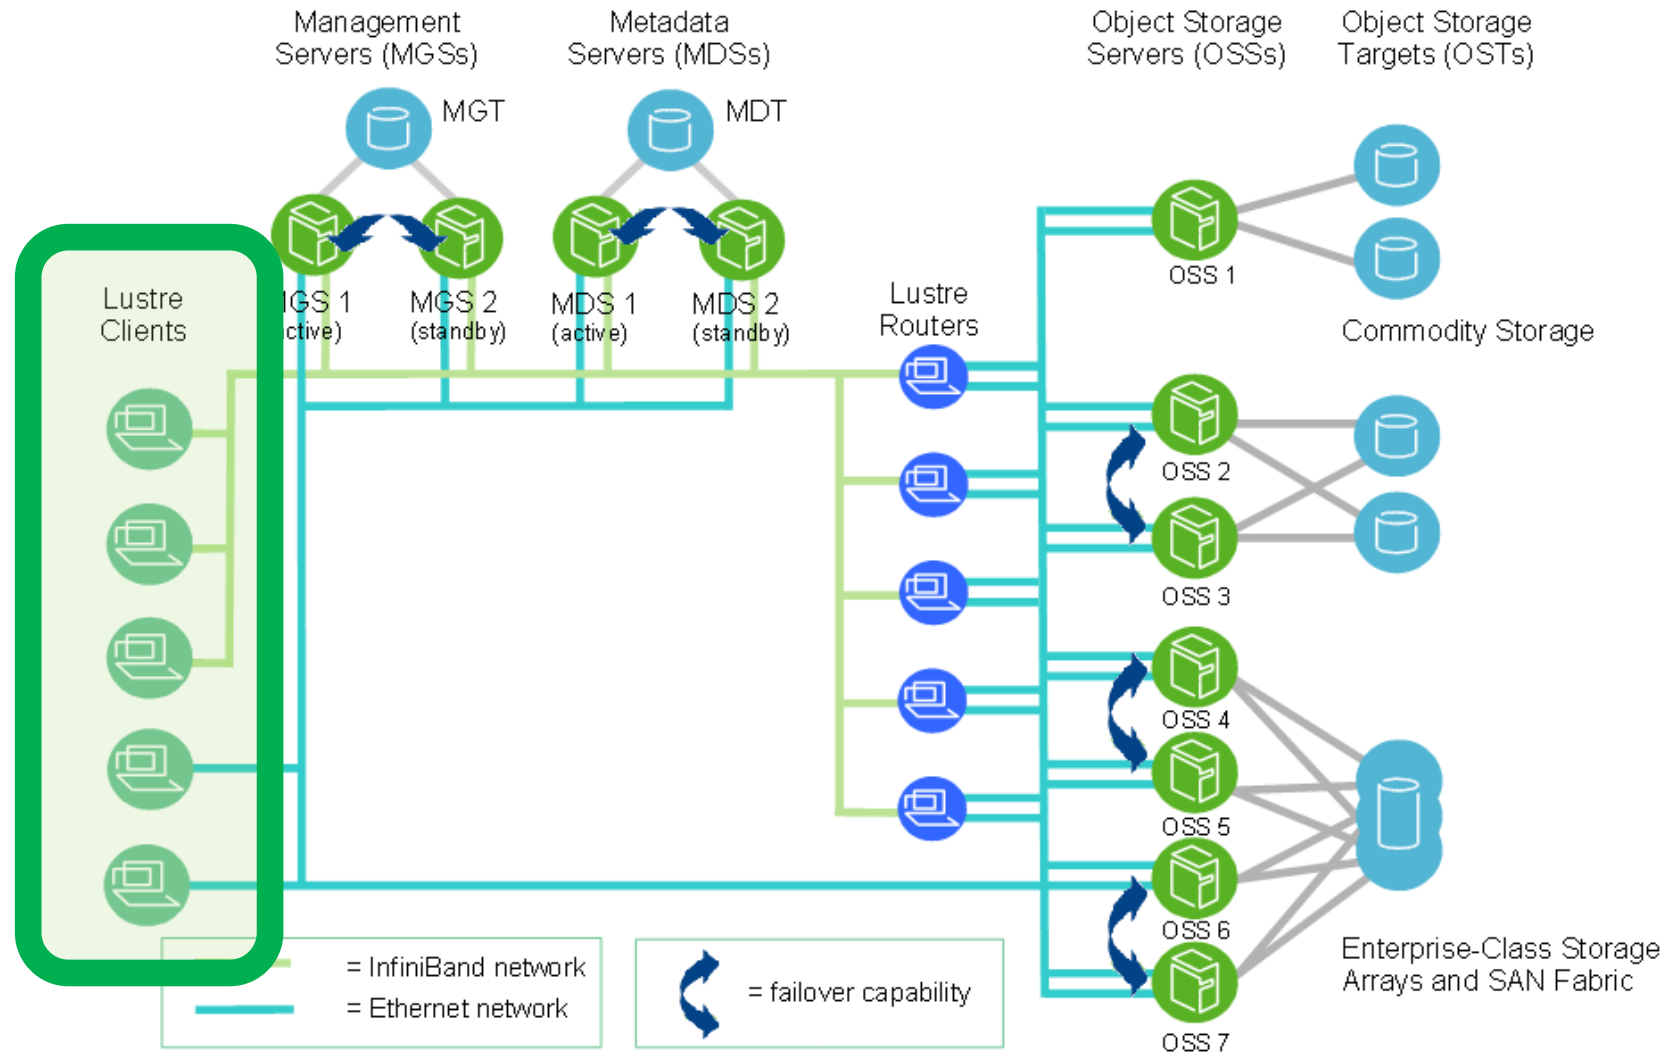
# Scratch – a Lustre file system



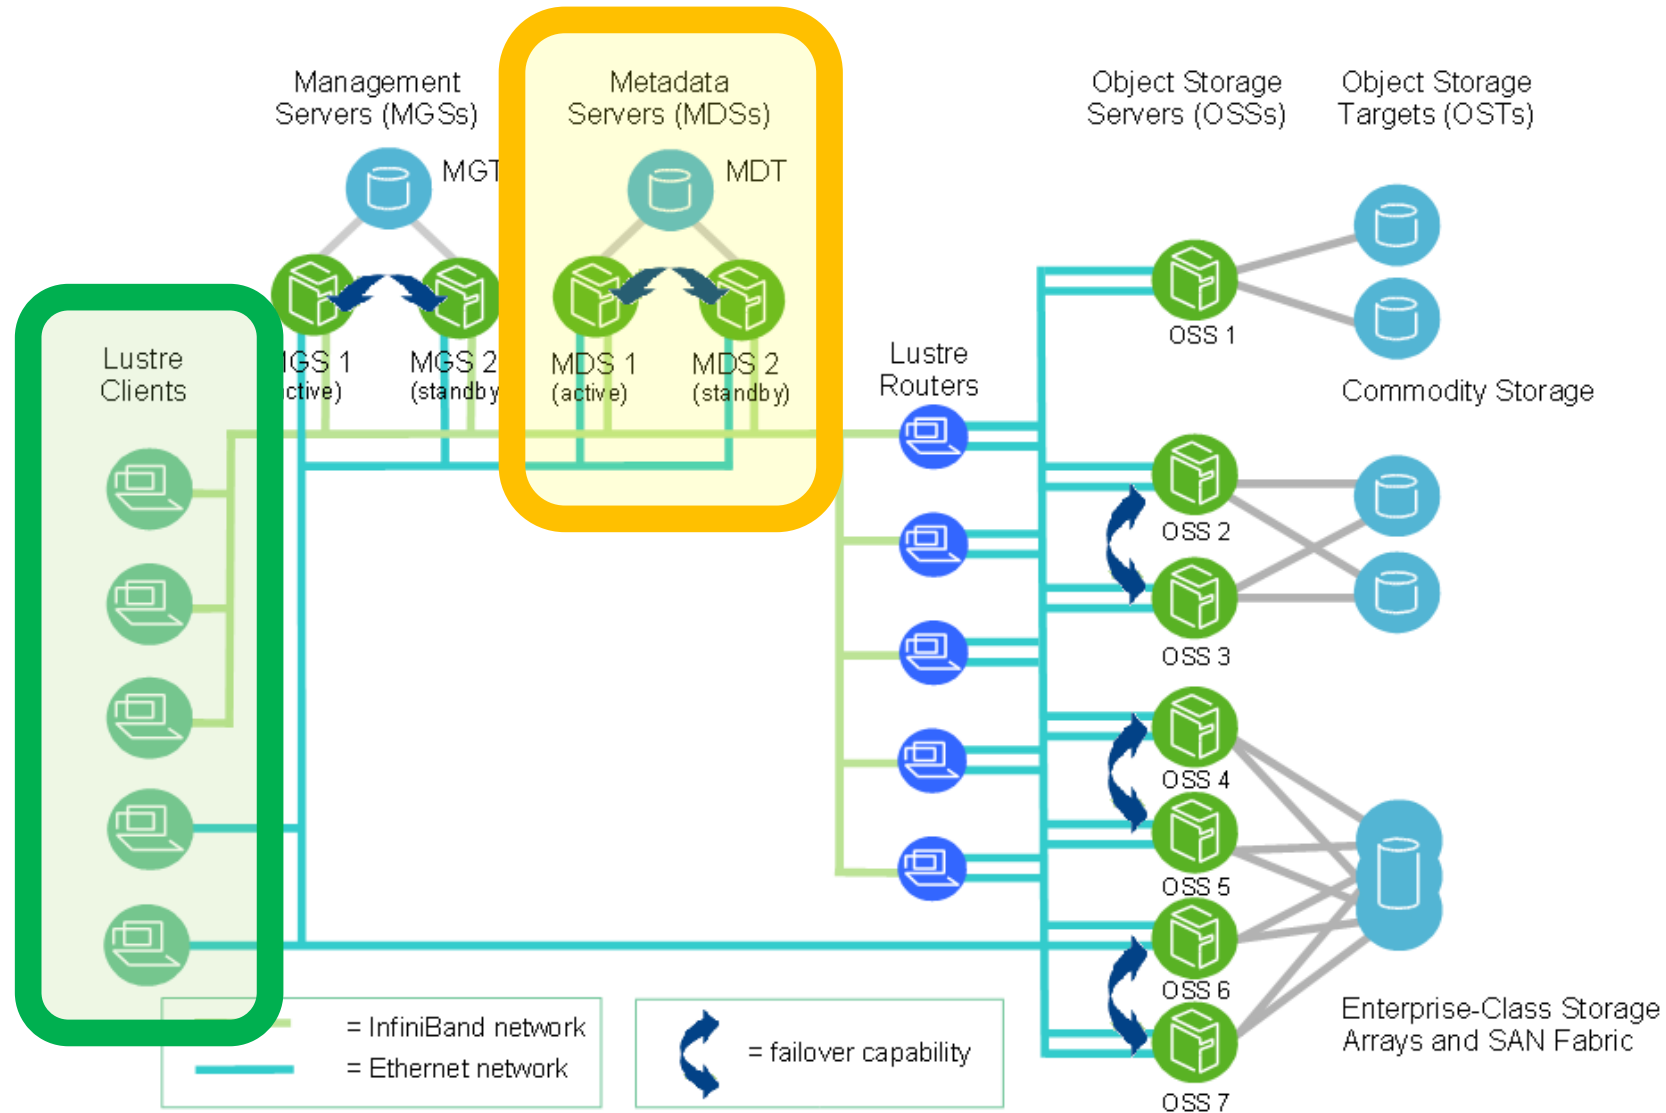
## Side note

It's not exactly Piz Daint's network layout: we have all InfiniBand + more routers

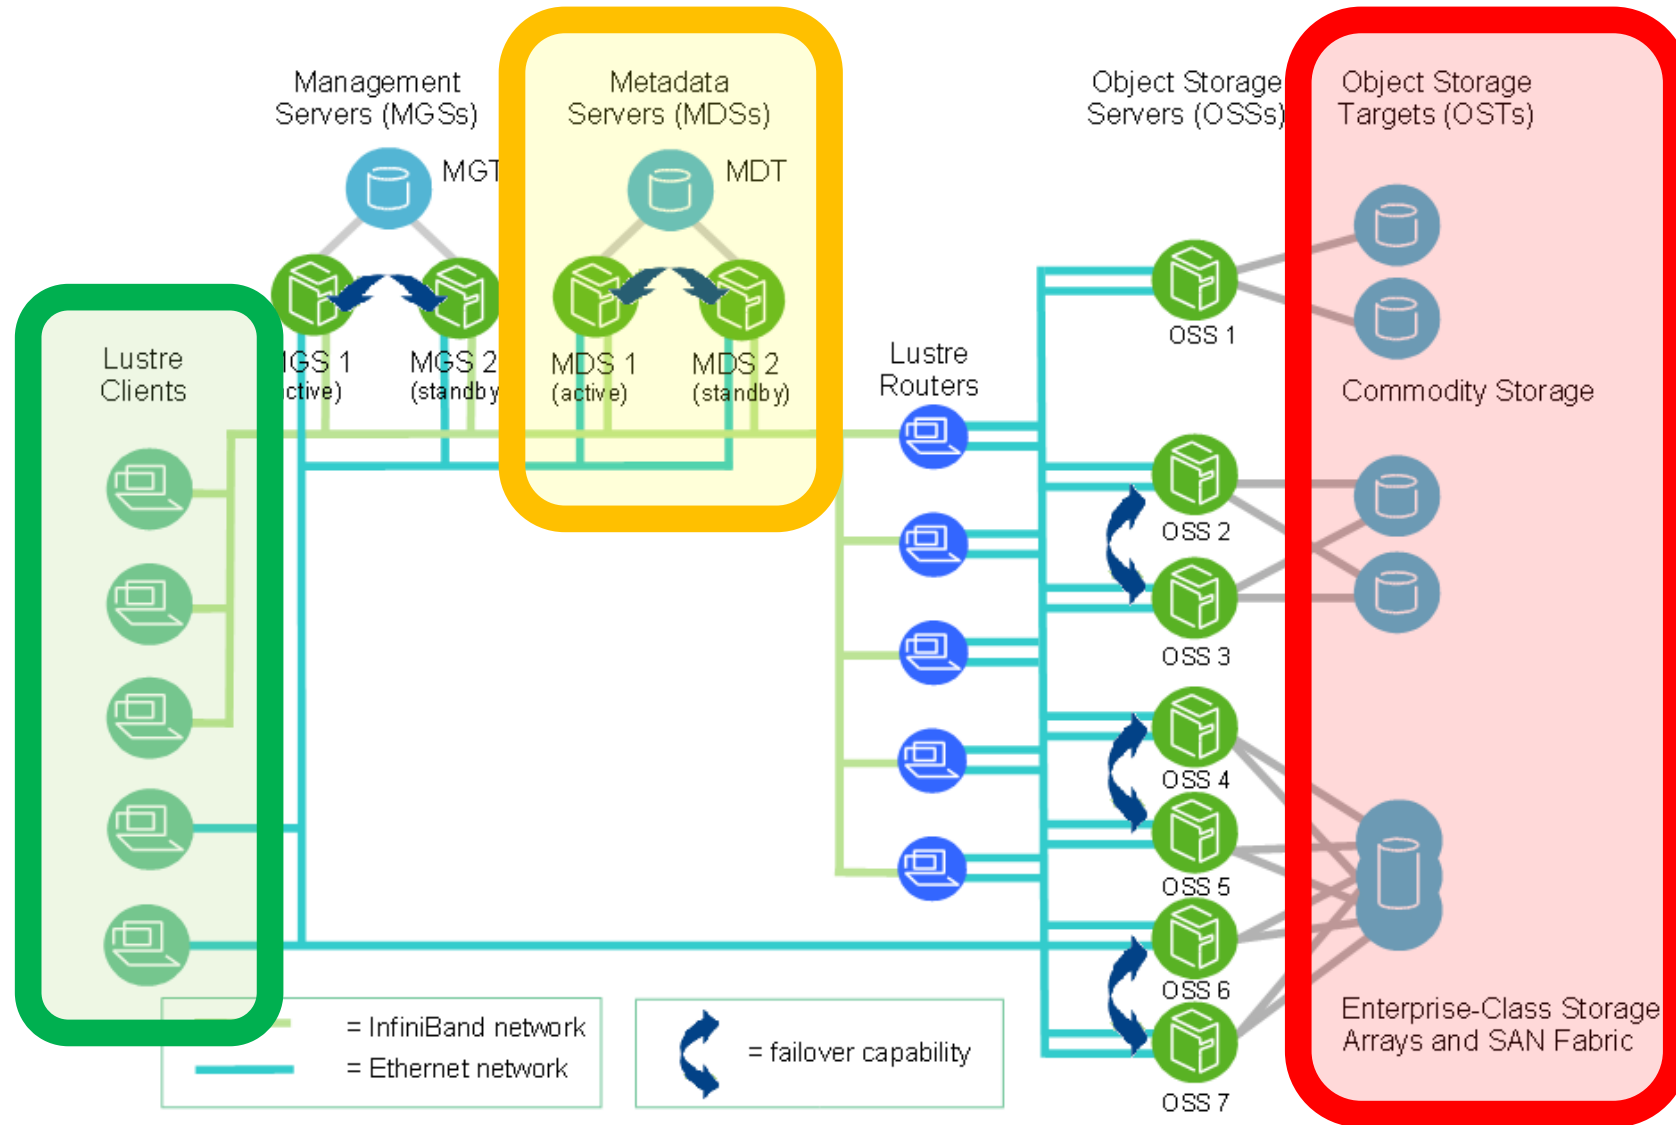
# Scratch – a Lustre file system



# Scratch – a Lustre file system

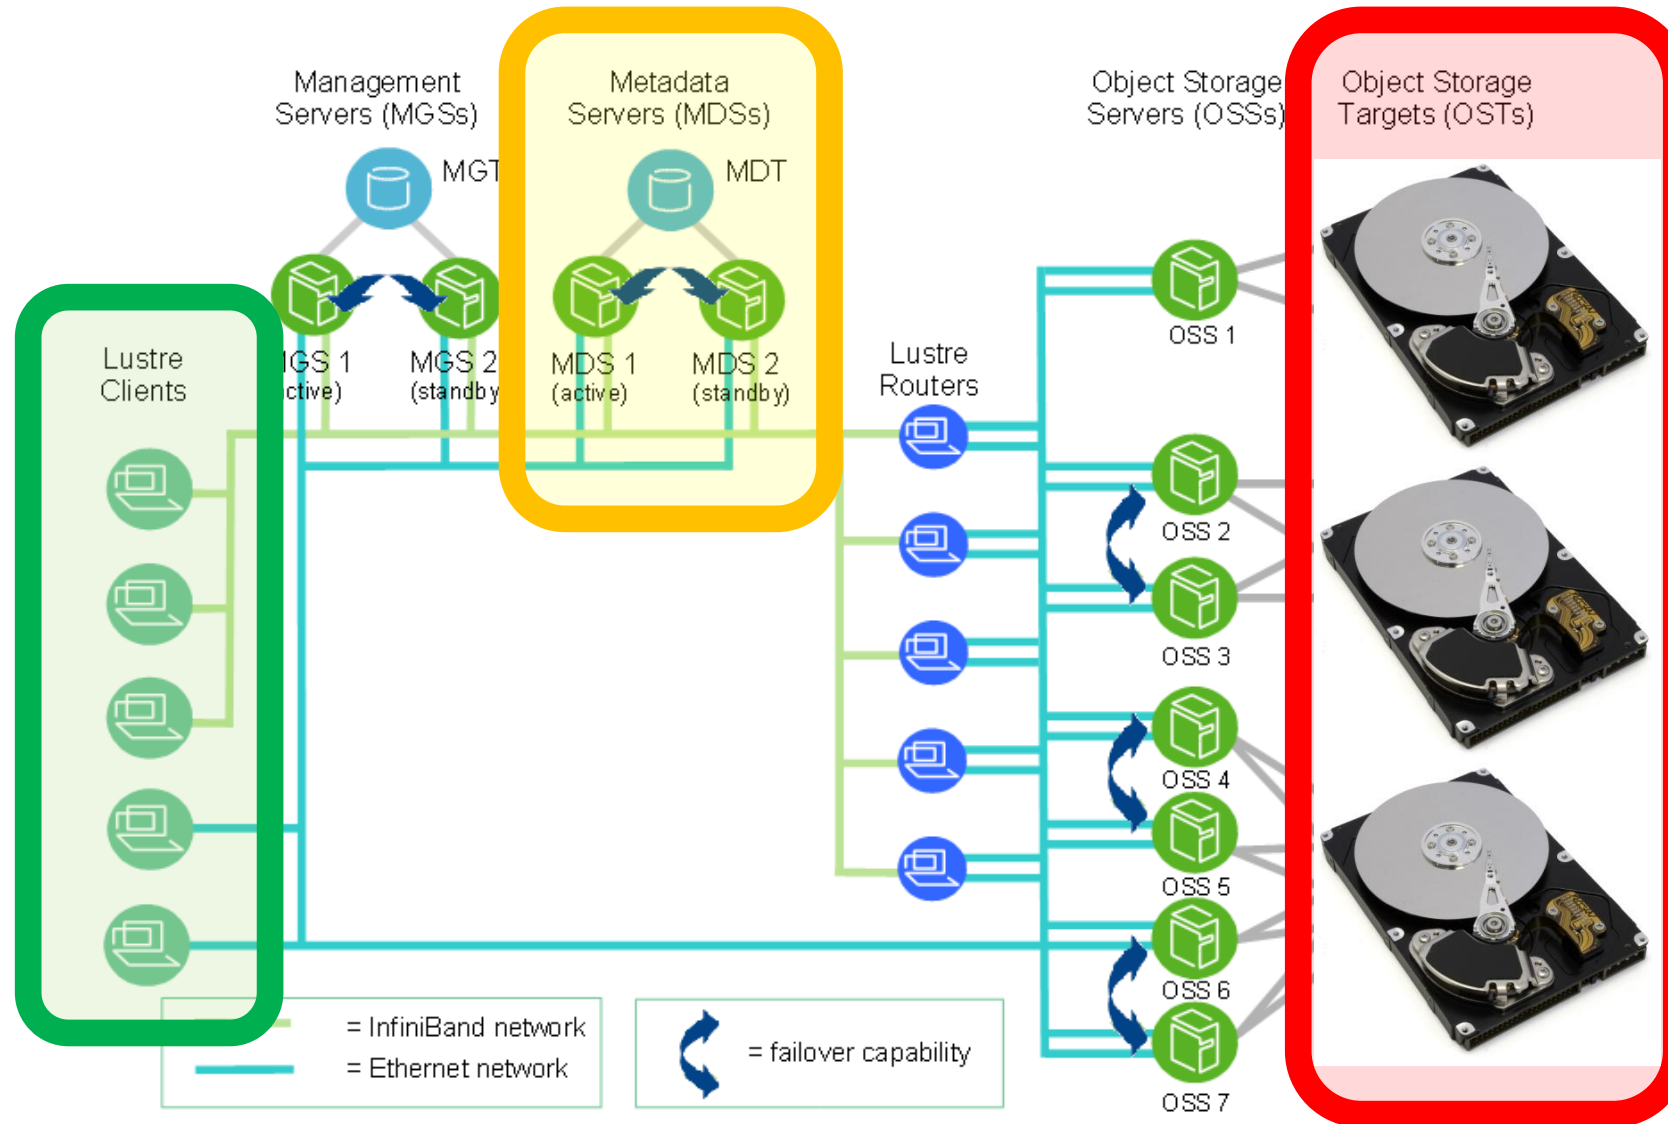


# Scratch – a Lustre file system





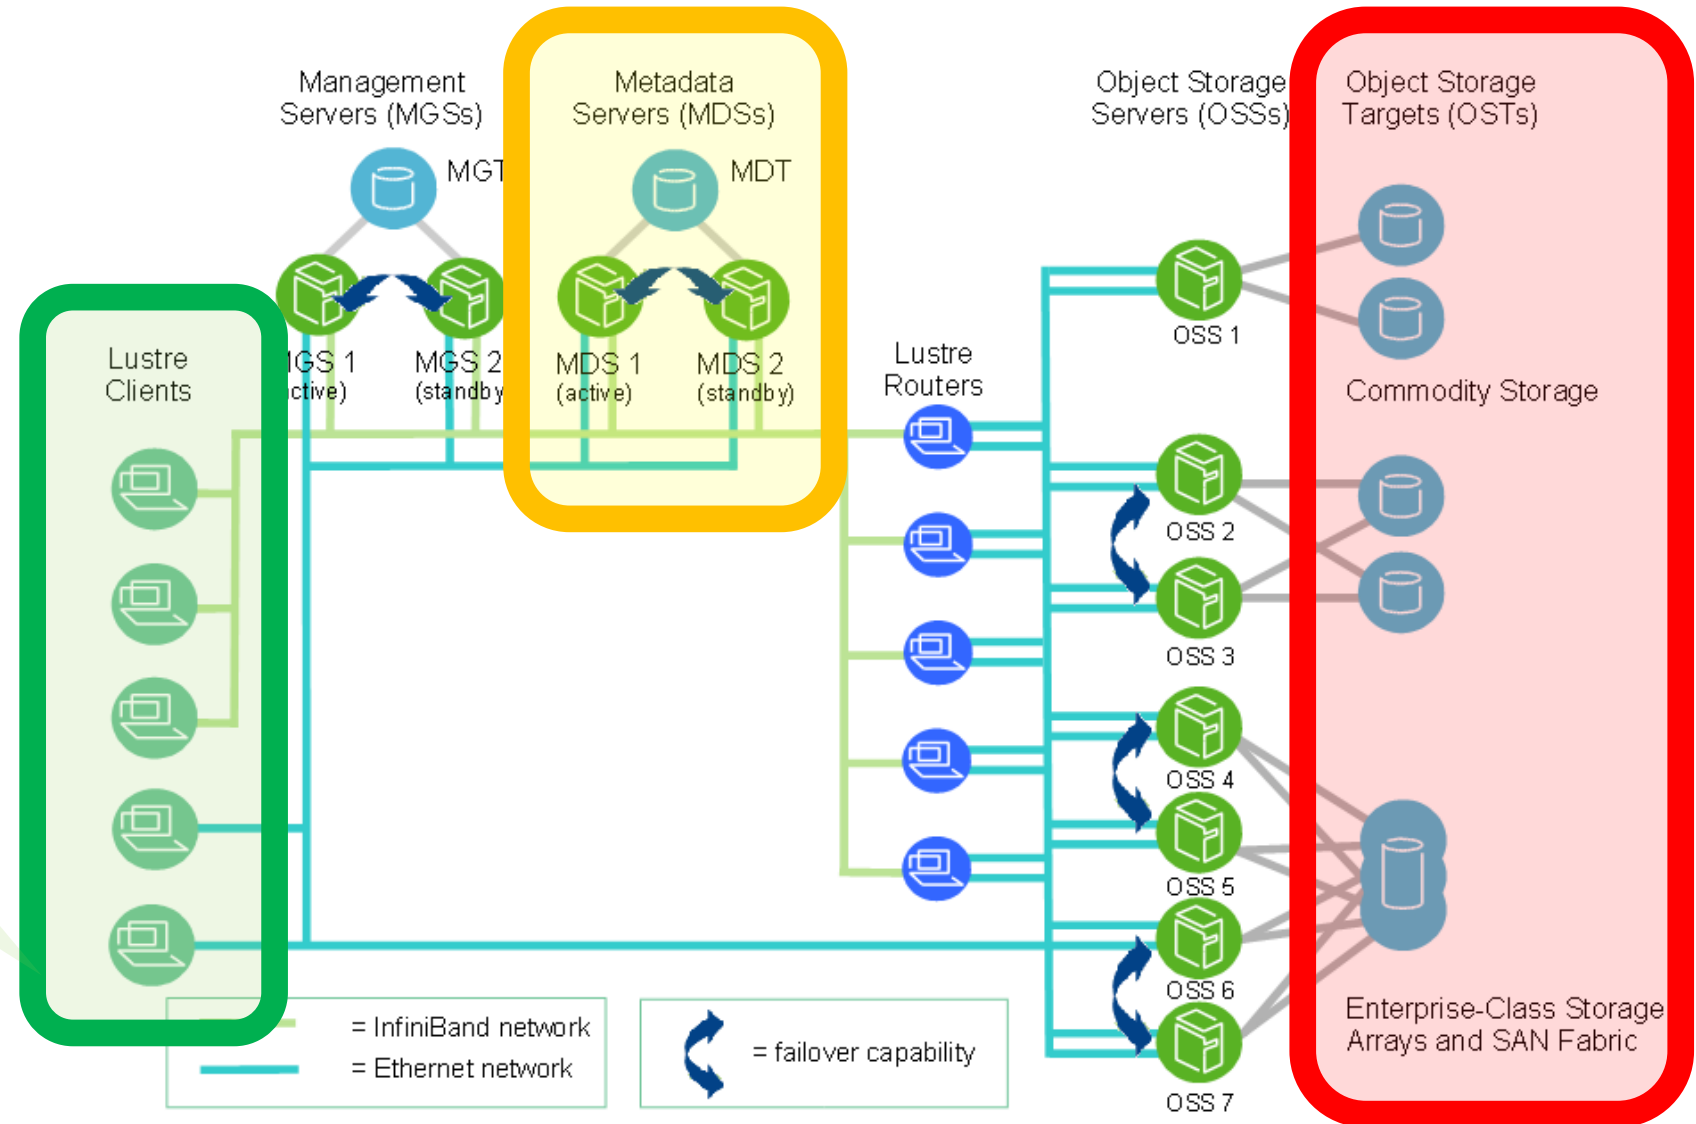
# Scratch – a Lustre file system



# Scratch – a Lustre file system

**\$SCRATCH:**  
/scratch/snx3000/\$USER

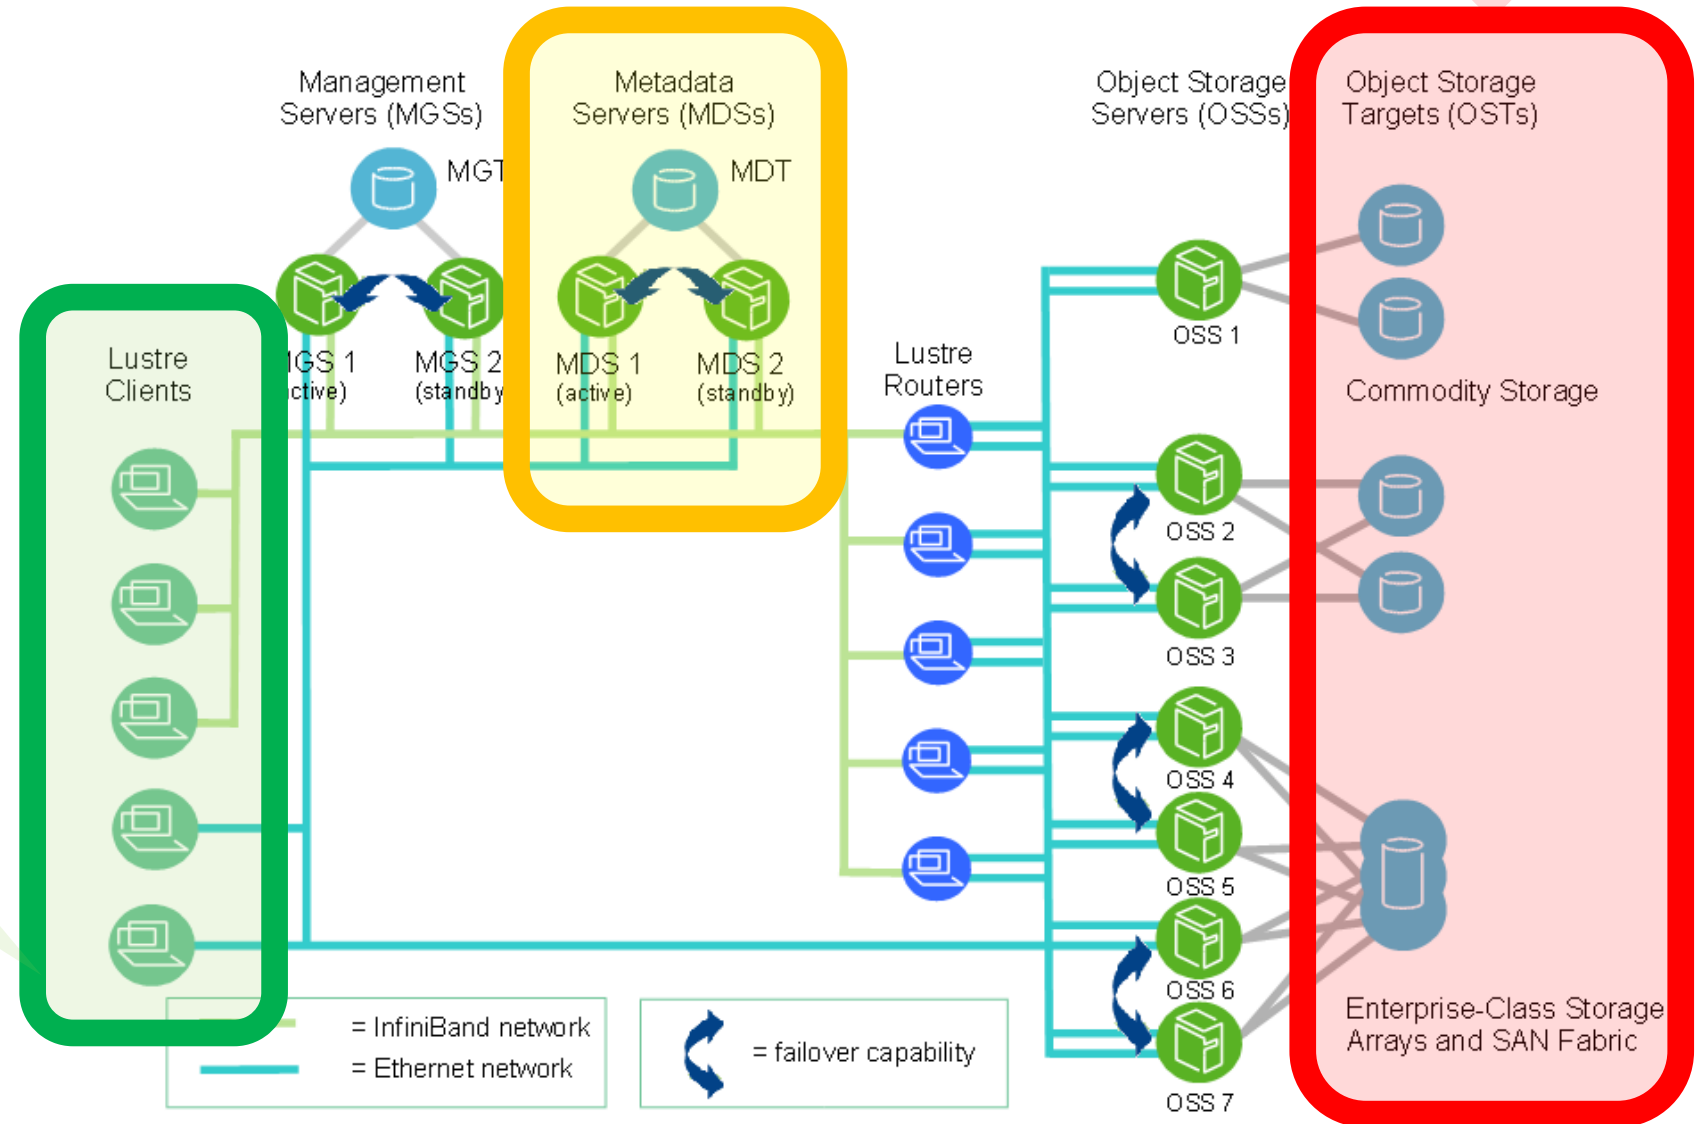
**snx3000:**  
Cray Sonexion 3000  
(>100 GB/s)



# Scratch – a Lustre file system

**\$SCRATCH:**  
/scratch/snx3000/\$USER

**snx3000:**  
Cray Sonexion 3000  
(>100 GB/s)

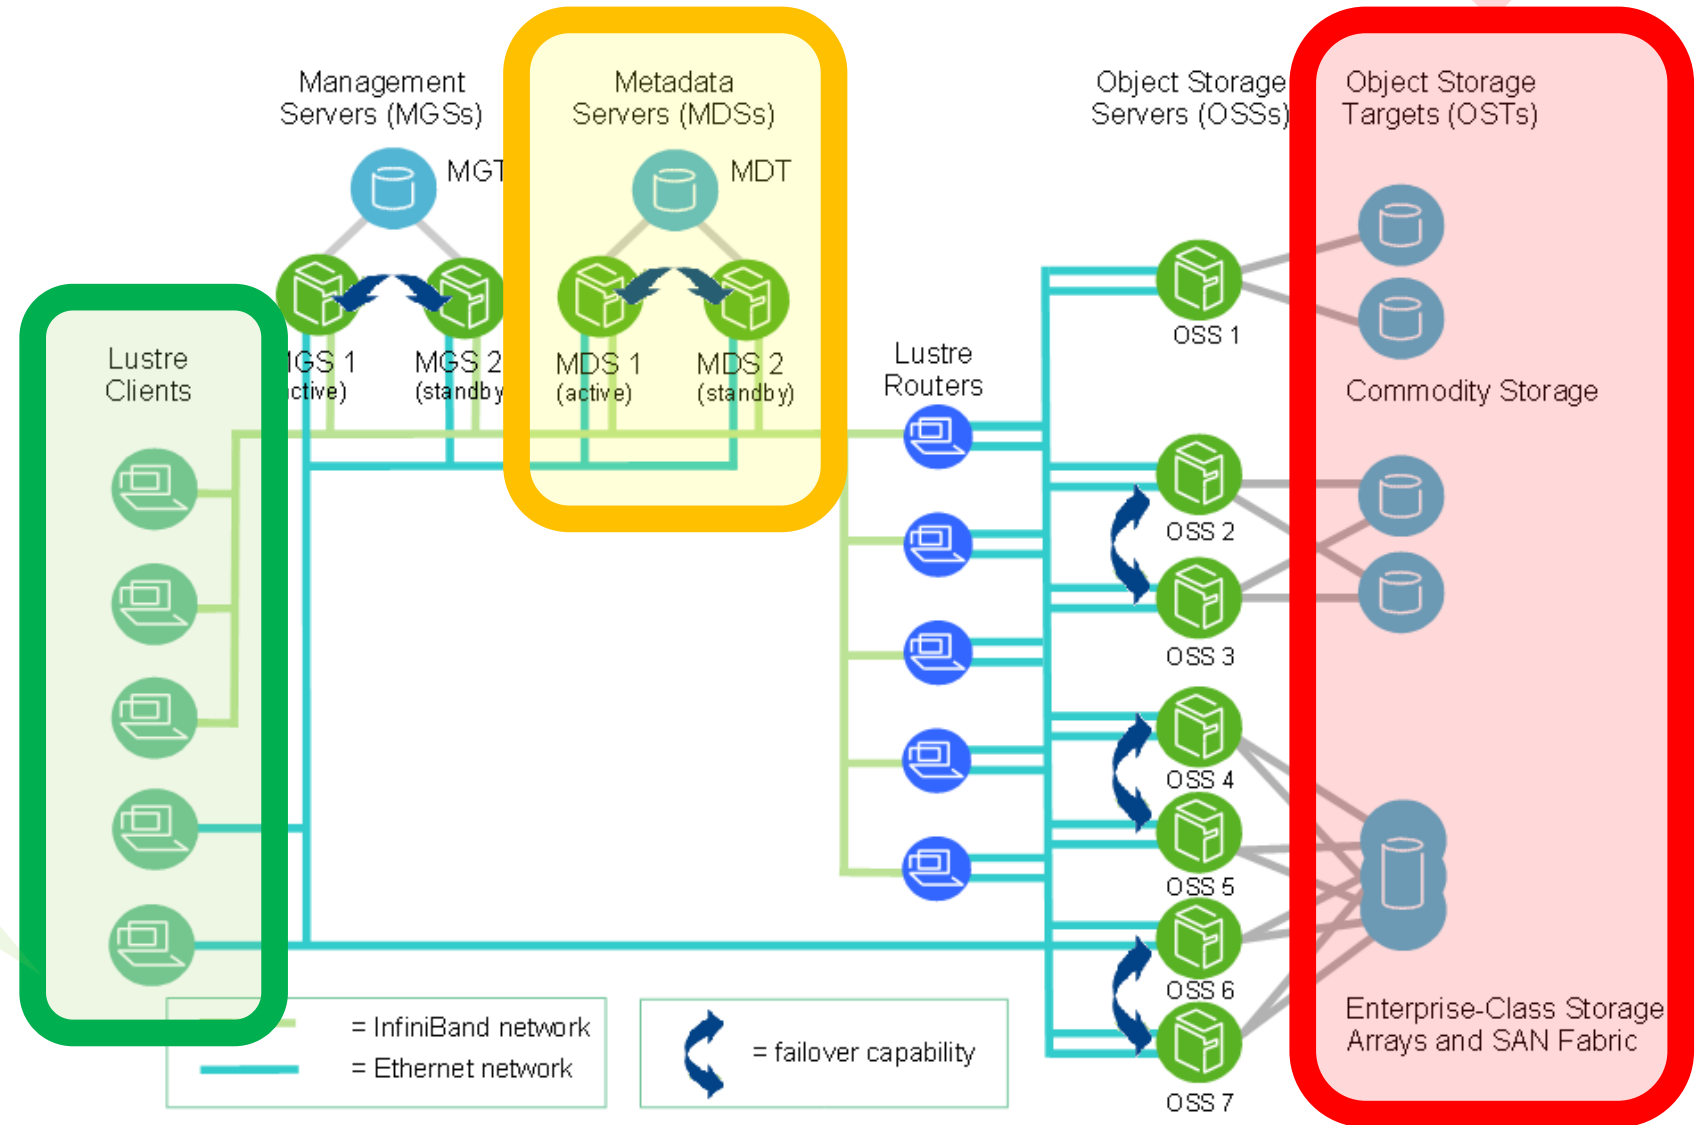


# Scratch – a Lustre file system

**\$SCRATCH:**  
/scratch/snx3000/\$USER

**snx3000:**  
Cray Sonexion 3000  
(**>100 GB/s**)

**Only possible if  
all or most OSTs  
accessed  
=> parallel I/O**





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre


**ETH** zürich

# Parallel I/O? – Common approaches


---

# Parallel I/O? – Common approaches

- File-per-process

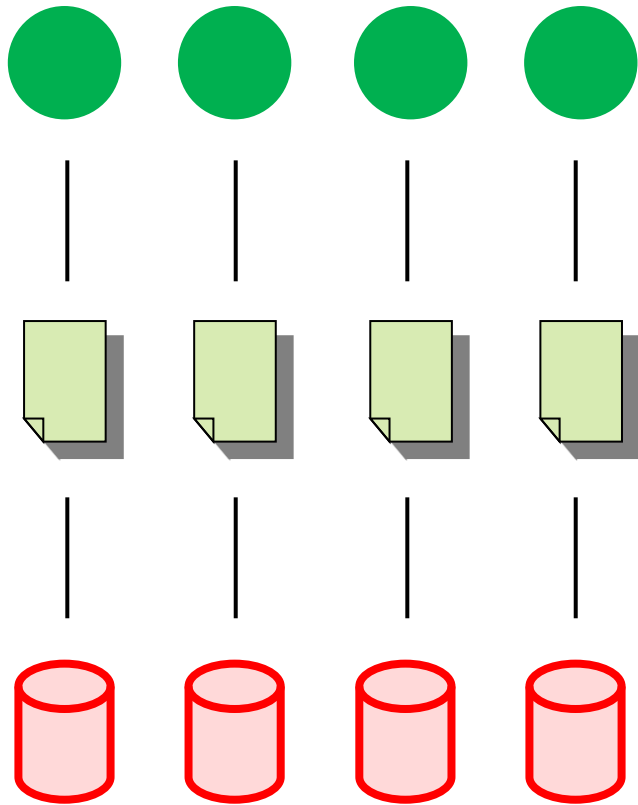


- Shared file

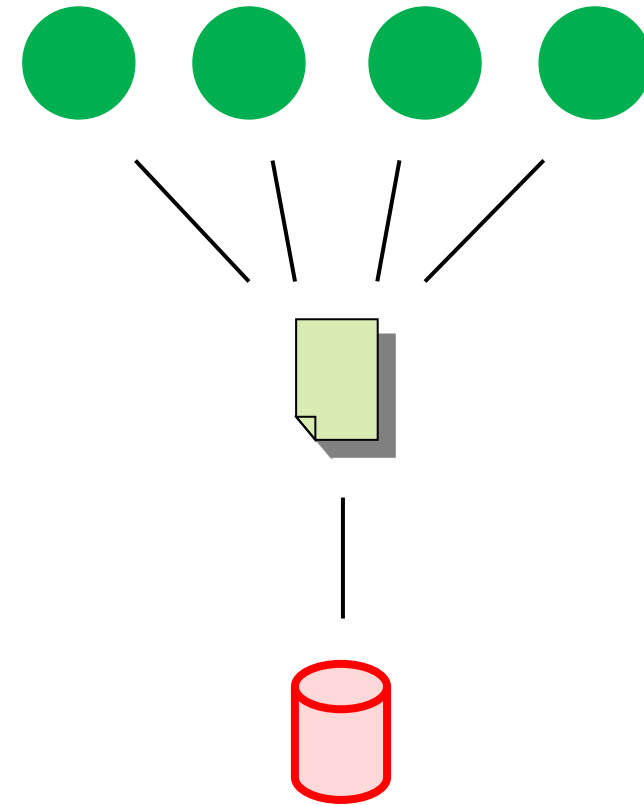


# Parallel I/O? – Common approaches

- File-per-process

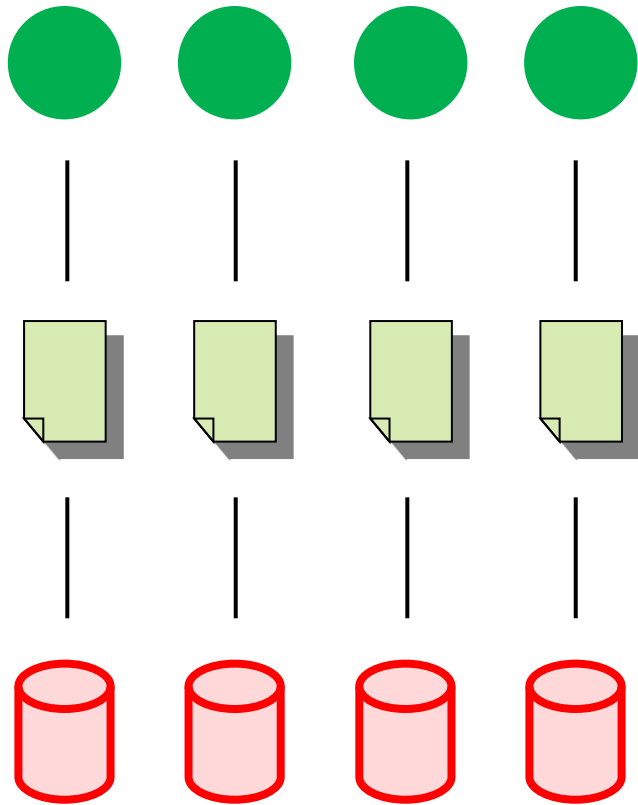


- Shared file




# Parallel I/O? – Common approaches

- File-per-process



- Shared file

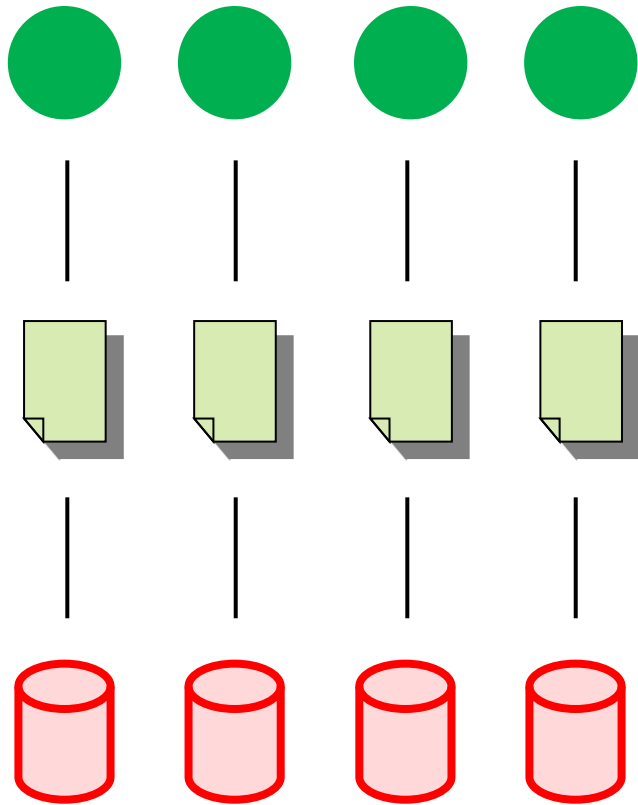


**Only  
1 OST  
accessed  
by default!**

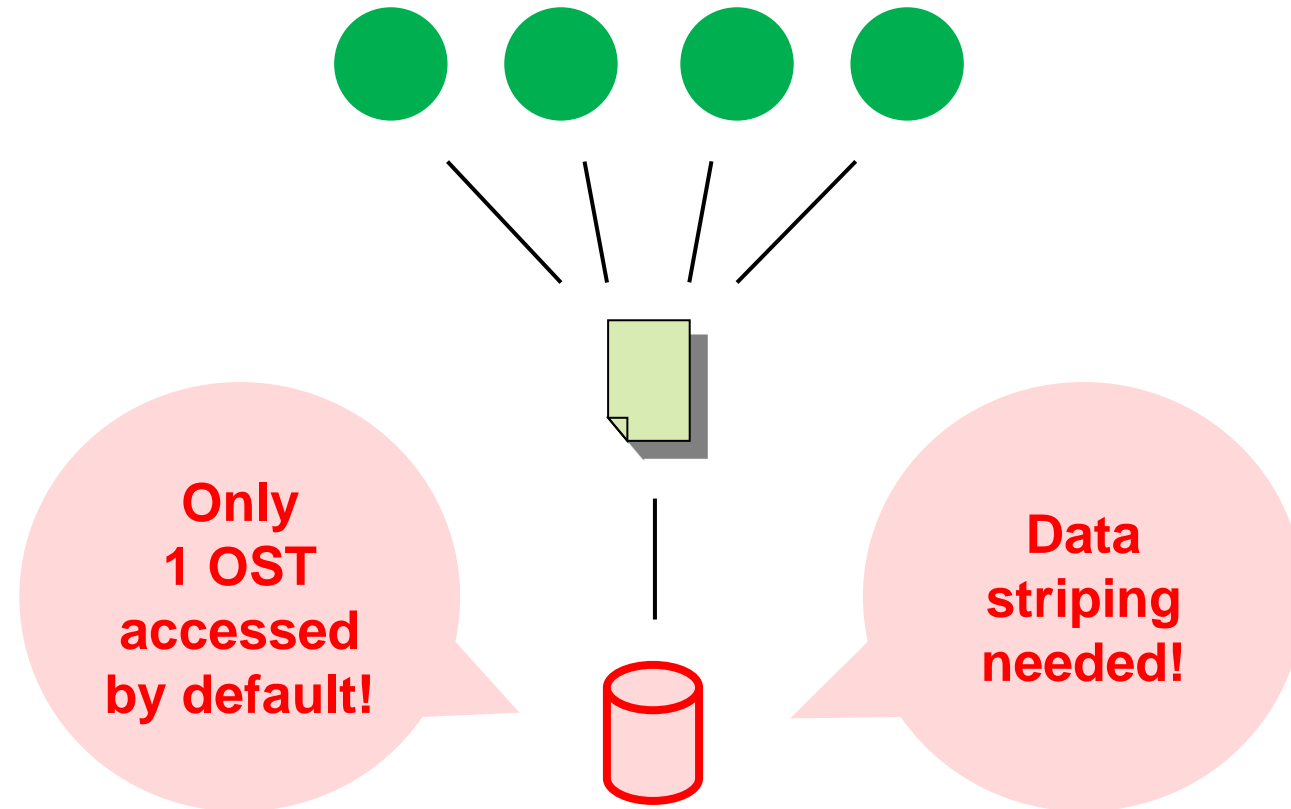


# Parallel I/O? – Common approaches

- File-per-process

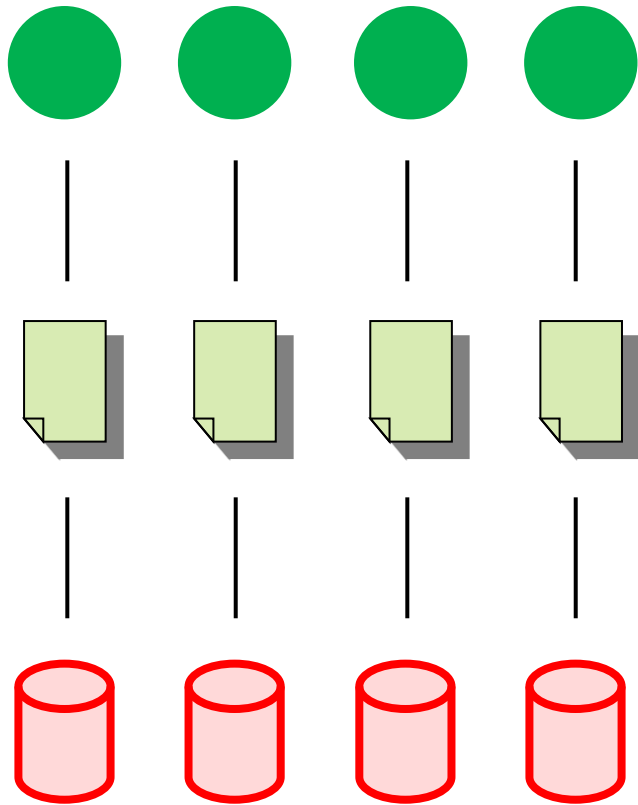


- Shared file

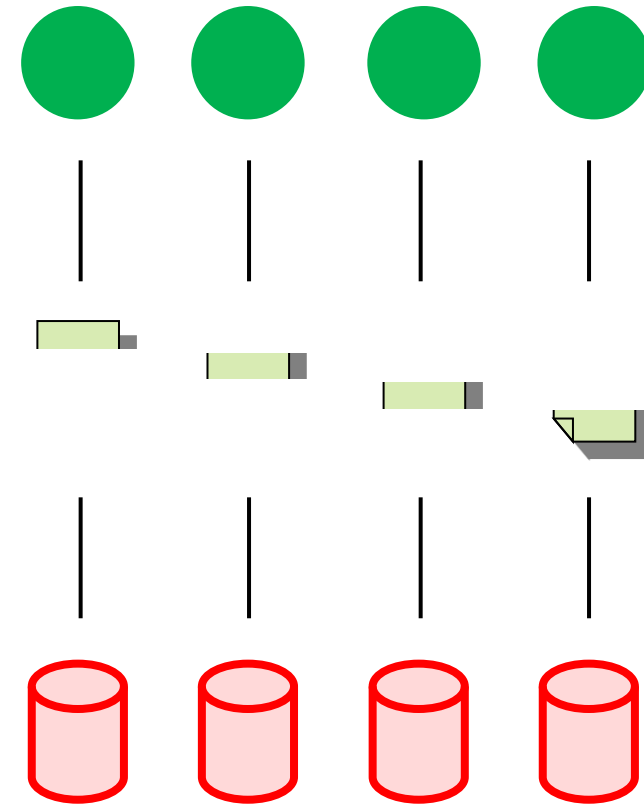


# Parallel I/O? – Common approaches

- File-per-process

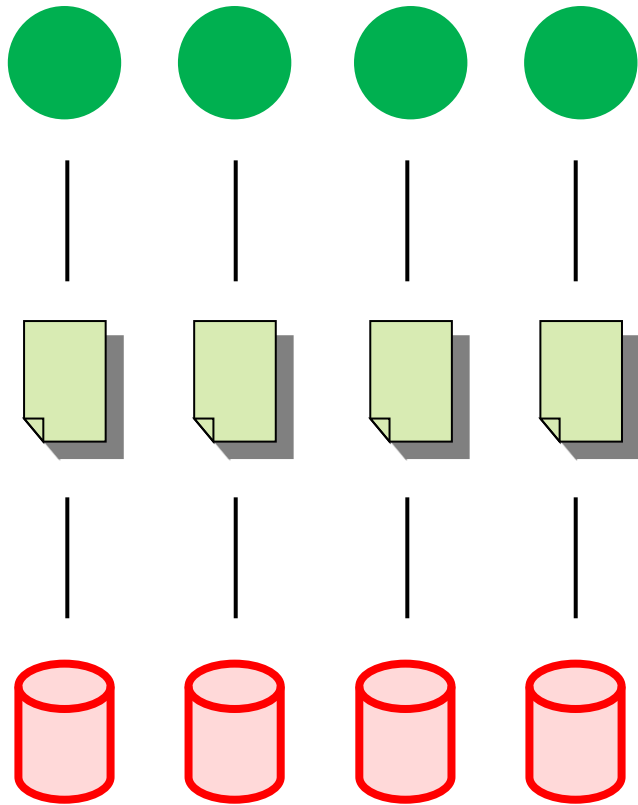


- Shared file (with striping)

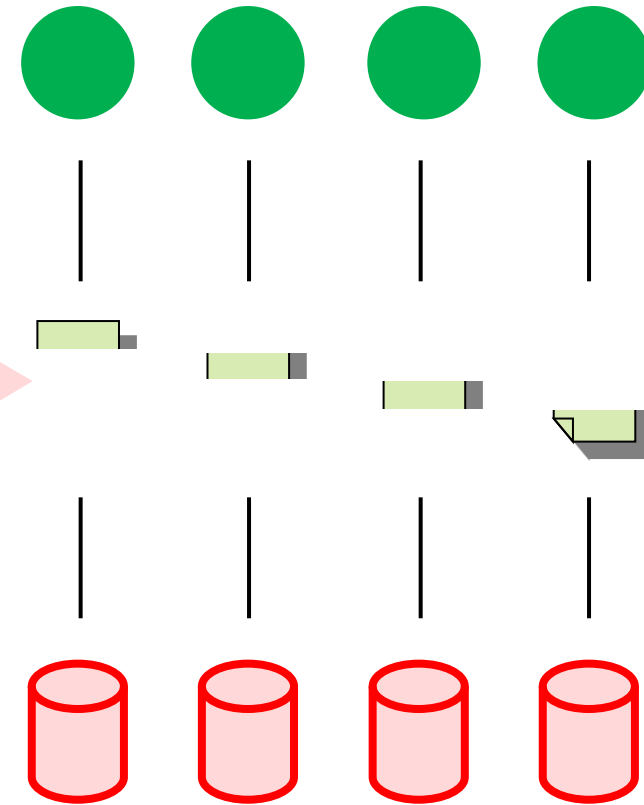


# Parallel I/O? – Common approaches

- File-per-process



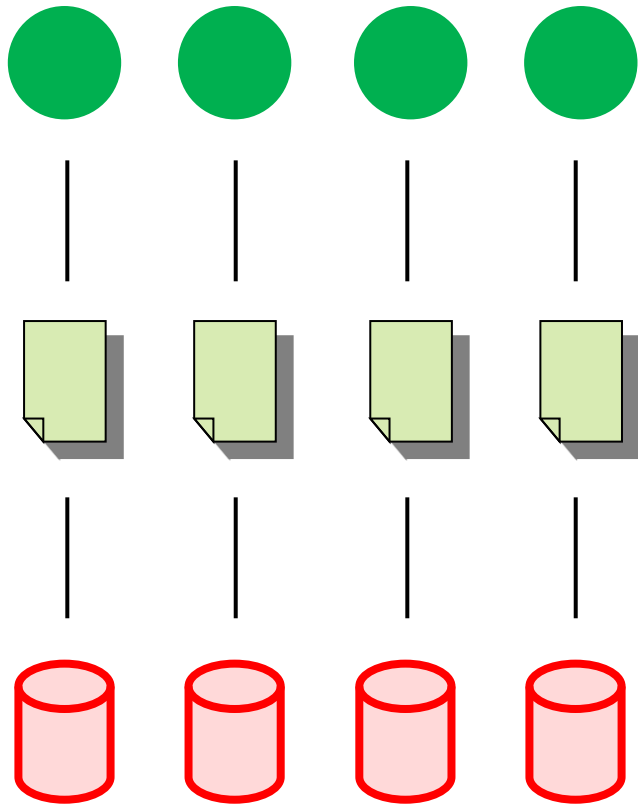
- Shared file (with striping)



**Stripe  
count = 4**  
(distribute  
file on 4  
OSTs)

# Parallel I/O? – Common approaches

- File-per-process



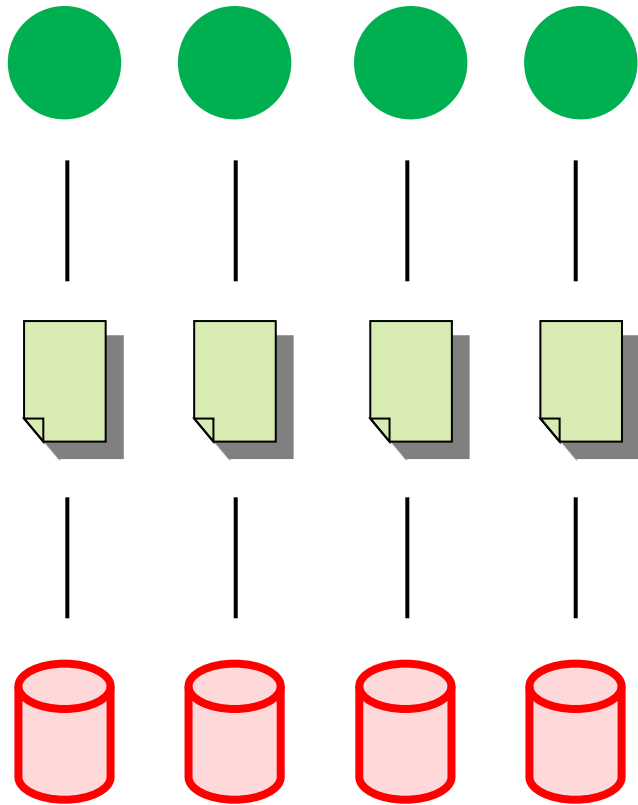
**Maximal simplicity is possible** on the application level (no I/O library is needed).



Creates **many files** if used at large scale (=> risk of contention of metadata servers).

# Parallel I/O? – Common approaches

- File-per-process



## General recommendations

- Set *stripe count* = 1 (default).
- Don't create thousand of files in a single folder (group files in subfolders).
- Don't access thousand of files simultaneously.
- Be aware that if you cause contention, all users will suffer (all I/O resources are shared).

# Parallel I/O? – Common approaches



Keeping the **number of files** small is straightforward.

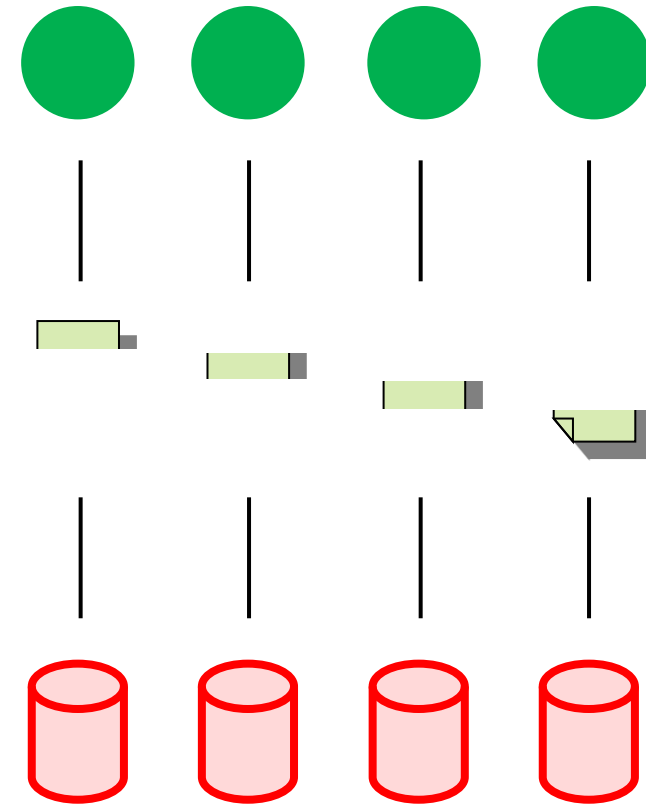


**Post-processing is simple** with common tools as I/O libraries make it easy to  
(1) add extensive metadata, and  
(2) create shared files as if it was done by a single process.



**Careful tuning** is often required to reach good I/O performance.

- Shared file (with striping)

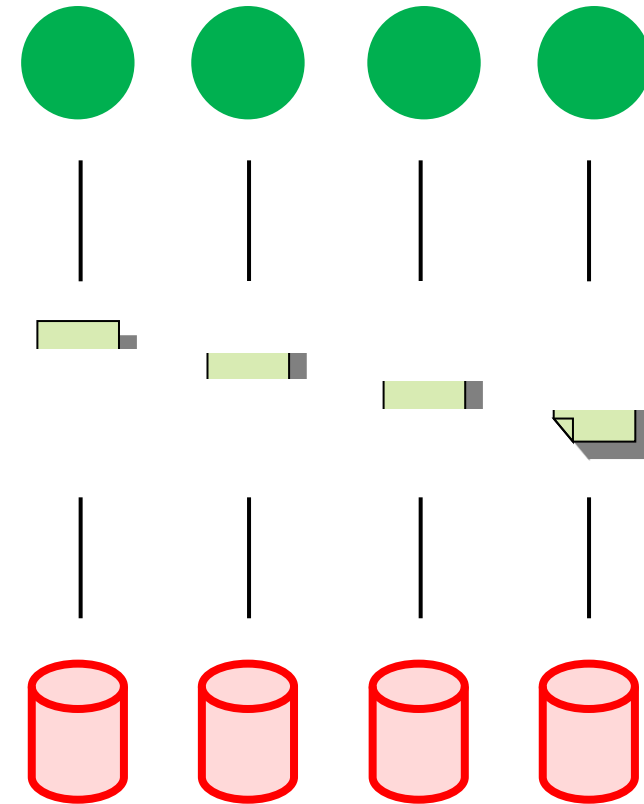


# Parallel I/O? – Common approaches

## General recommendations (1/5)

- If you access **one large shared file** (GBs), set *stripe count* = *#OSTs*.
- If you simultaneously access **multiple large shared files**, set *stripe count* such that  $\#files * \textit{stripe count} = k * \#OSTs$ , where  $k \in \{1, \dots, 4\}$ .
- *#OSTs* = 40 on snx3000. Nevertheless, try the above formulas also with the value 32 (better alignment possible).
- For any small file, set *stripe count* = 1.

- Shared file (with striping)

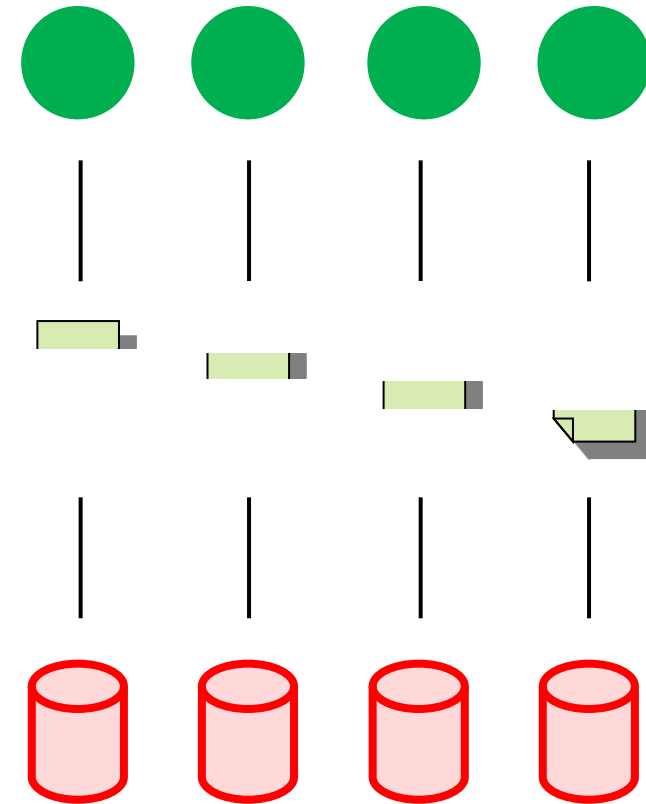


# Parallel I/O? – Common approaches

## General recommendations (2/5)

- Convenient: set the striping configuration for your simulation output folder(s), e.g.  
`lfs setstripe --stripe-count 32 <output folder>`  
=> Inside, files will be created with the same striping configuration as the folder(s) itself.  
**Don't copy the executable in there!**
- Check the striping configuration of a file or folder:  
`lfs getstripe <file/folder>`
- More information: `lfs --help`

- Shared file (with striping)





# Parallel I/O? – Common approaches

## General recommendations (3/5)

- Use **collective I/O operations** (enable merging of I/O requests of different processes into fewer larger ones).

MPIIO:

use functions with suffix ‘\_all’  
(e.g. `MPI_File_write_all`)


HDF5:

`H5Pset_dxpl_mpio(..., H5FD_MPIO_COLLECTIVE);`

NetCDF:

`nc_var_par_access(..., NC_COLLECTIVE);`

- Shared file (with striping + collective buffering)



# Parallel I/O? – Common approaches

## General recommendations (3/5)

- Use **collective I/O operations** (enable merging of I/O requests of different processes into fewer larger ones).

MPIIO:

use functions with suffix ‘\_all’  
(e.g. `MPI_File_write_all`)

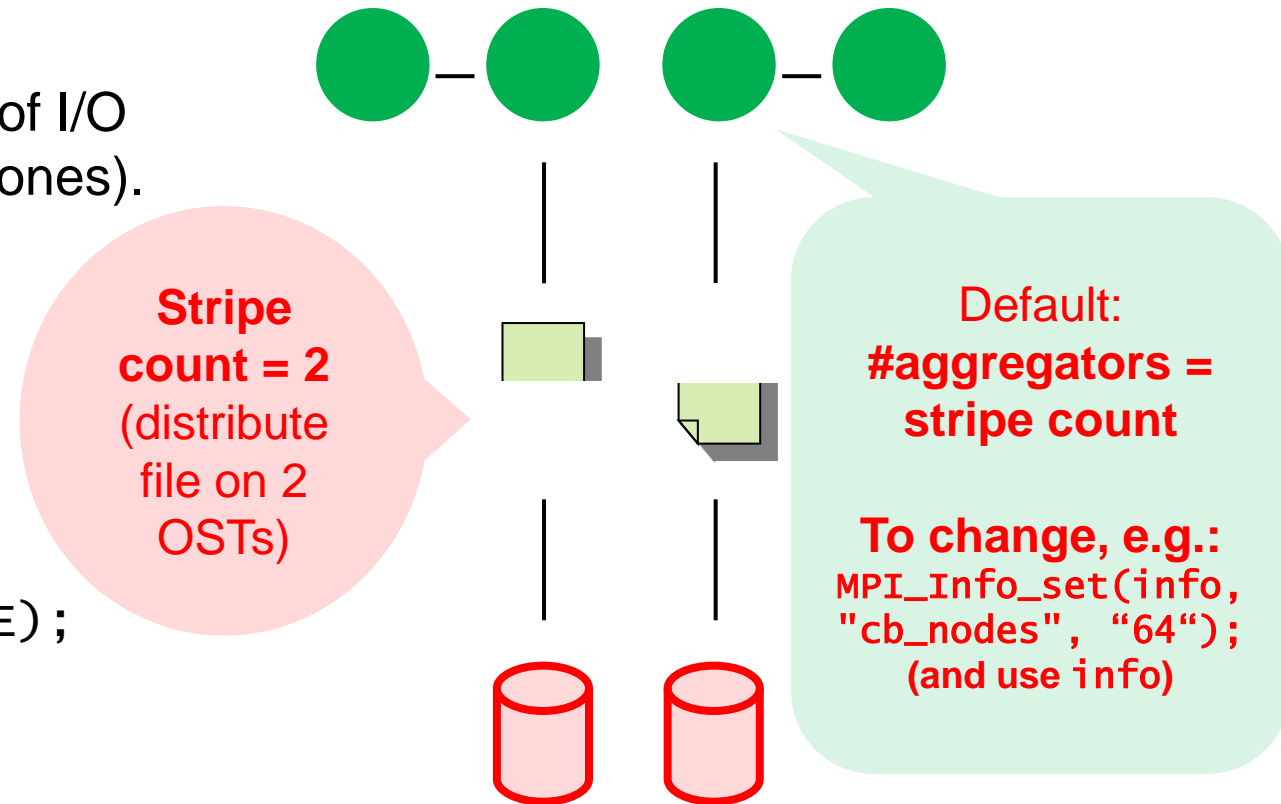
HDF5:

```
H5Pset_dxpl_mpio(..., H5FD_MPIO_COLLECTIVE);
```

NetCDF:

```
nc_var_par_access(..., NC_COLLECTIVE);
```

- Shared file (with striping + collective buffering)

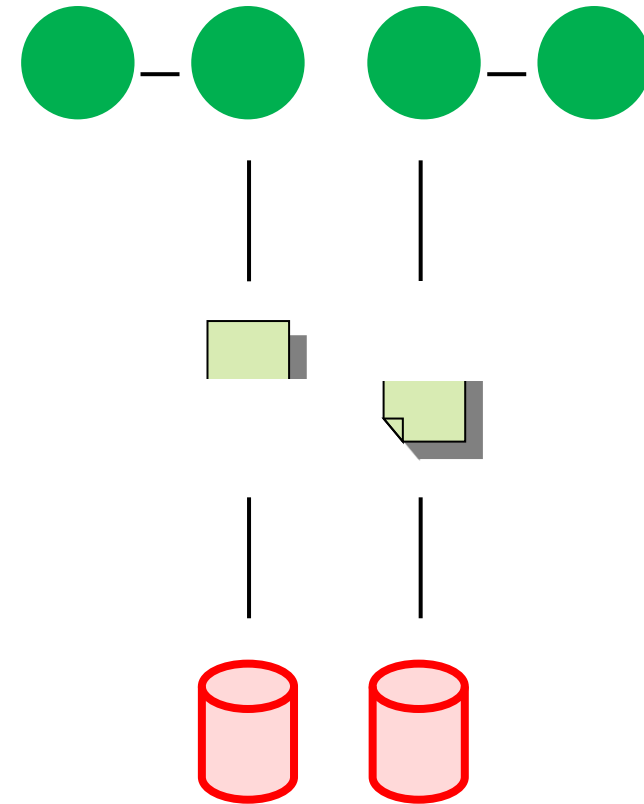


# Parallel I/O? – Common approaches

## General recommendations (4/5)

- Print used MPIIO hints (e.g. `cb_nodes`; also used for HDF5 and NetCDF as have MPIIO underneath!):  
`export MPICH_MPIIO_HINTS_DISPLAY=1`
- Print statistics about I/O:  
`export MPICH_MPIIO_STATS=1`
- Look up detailed information on MPIIO hints:  
`man intro_mpi`

- Shared file (with striping + collective buffering)




# Parallel I/O? – Common approaches

## General recommendations (5/5)

- Refer to the **websites of the parallel I/O libraries** for details on their best usage, e.g.:  
*HDF5*: [www.hdfgroup.org](http://www.hdfgroup.org)  
*NetCDF*: [www.unidata.ucar.edu/software/netcdf/](http://www.unidata.ucar.edu/software/netcdf/)  
*ADIOS*: [www.olcf.ornl.gov/center-projects/adios/](http://www.olcf.ornl.gov/center-projects/adios/)
- Follow widely adopted **Metadata conventions** to enable straightforward pre- and post-processing. E.g:  
*NetCDF CF Metadata Conventions*: [cfconventions.org](http://cfconventions.org)  
*XDMF*: [www.xdmf.org](http://www.xdmf.org)  
*GADGET*: [wwwmpa.mpa-garching.mpg.de/gadget](http://wwwmpa.mpa-garching.mpg.de/gadget)

- Shared file (with striping + collective buffering)




# Parallel I/O? – Common approaches

## Available parallel I/O libraries on Piz Daint

- `cray-hdf5-parallel` (HDF5)
- `cray-netcdf-hdf5parallel` (NetCDF using HDF5 underneath)
- `cray-parallel-netcdf` (NetCDF using PnetCDF underneath)

- Shared file (with striping + collective buffering)





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# General recommendations for any I/O approach on Piz Daint

---

# General recommendations for any I/O approach on Piz Daint

- **No ASCII**, except for small parameter files (easily 10 - 100 times slower than binary)
- Avoid opening and closing files frequently.
- Do not open files for read and write access, but instead for read-only or write-only.
- **Avoid small and frequent I/O request.**
- Avoid random file access; regular access patterns work best in general.
- Avoid multiple processes accessing the same data.
- Read small files just from one process and broadcast the data to the remaining.
- **Limit file metadata access as much as possible on the Lustre file system**; in particular, avoid the usage of 'ls -l' and use instead 'ls' or 'lsfs find' when possible.
- **Be aware that if you cause contention on the file system, all users will suffer the slowdowns as all I/O resources are shared.**

# Conclusions

---



# Conclusions

- *parallel I/O lib*  $\neq$  *good performance*
  - $\Rightarrow$  Careful tuning is often required (striping + collective optimizations)
- **Metadata conventions** enable straightforward pre- and post-processing and data portability between applications in general.



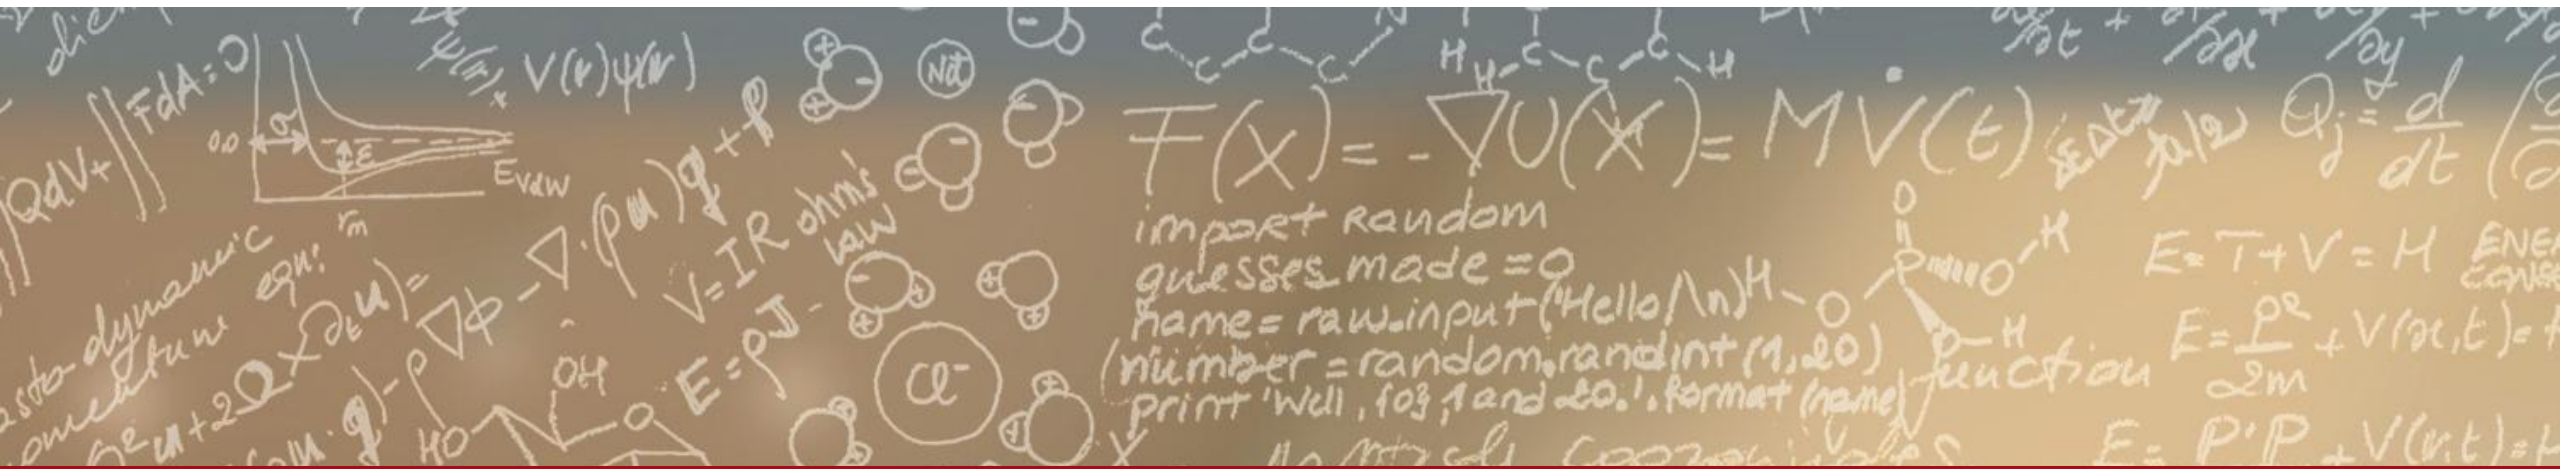
*Piz Daint in the machine room at CSCS*



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your kind attention**